

FDS Step 4

Step 1: Derive sentiment distribution across the document

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.decomposition import LatentDirichletAllocation
from sklearn.feature_extraction.text import CountVectorizer
from statsmodels.tsa.seasonal import seasonal_decompose
from wordcloud import WordCloud

# Load your CSV file
df = pd.read_csv('path_to_your_file.csv')

# Basic Preprocessing
# Ensure that 'timestamp' is in datetime format if it exists
# df['timestamp'] = pd.to_datetime(df['timestamp'])

# Create additional features (e.g., comment length)
df['comment_length'] = df['comment'].apply(len)

# --- Time Series Analysis ---

# If you have timestamp data, you can decompose the sentiment time series
if 'timestamp' in df.columns:
    df.set_index('timestamp', inplace=True)
    result = seasonal_decompose(df['sentiment'], model='additive', period=30)
# Adjust period based on your data
    result.plot()
```

```

plt.show()

# --- Sentiment Distribution Analysis ---
sns.histplot(df['sentiment'], kde=True)
plt.title("Distribution of Sentiment Scores")
plt.xlabel("Sentiment Score")
plt.ylabel("Frequency")
plt.show()

```

Step 2: Generate wordclouds

```

# --- Word Cloud for Comment Analysis ---
text = " ".join(comment for comment in df['comment'])
wordcloud = WordCloud(max_words=100, background_color="white").generate(text)
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis("off")
plt.show()

```

Step 3: Perform predictive analytics

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb

# Load your CSV file
df = pd.read_csv('path_to_your_file.csv')

```

```

# Basic Preprocessing

# Ensure that 'timestamp' is in datetime format if it exists
# df['timestamp'] = pd.to_datetime(df['timestamp'])

# Create additional features (e.g., comment length)
df['comment_length'] = df['comment'].apply(len)

# Select features and target
X = df[['comment_length']] # Add other relevant features if available
y = df['sentiment']

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Initialize models
models = {
    'Linear Regression': LinearRegression(),
    'Decision Tree': DecisionTreeRegressor(random_state=42),
    'Support Vector Machine': SVR(),
    'Random Forest': RandomForestRegressor(random_state=42),
    'XGBoost': xgb.XGBRegressor(objective='reg:squarederror',
random_state=42)
}

# Train, predict, and evaluate each model
results = {}
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    results[name] = mse
    print(f'{name} - Mean Squared Error: {mse:.4f}')

```

```
# Compare the results

results_df = pd.DataFrame(list(results.items()), columns=['Model', 'Mean
Squared Error'])

sns.barplot(x='Mean Squared Error', y='Model', data=results_df)

plt.title("Model Comparison based on Mean Squared Error")

plt.show()
```

Step 4: Derive the plots

- a) **Residual plot:** A residuals plot shows the difference between the actual and predicted values. It helps to assess the model's fit.

```
# Plot residuals for the best-performing model
best_model_name = min(results, key=results.get)
best_model = models[best_model_name]

# Predict and calculate residuals
y_pred = best_model.predict(X_test)
residuals = y_test - y_pred

plt.figure(figsize=(10, 6))
plt.scatter(y_pred, residuals, alpha=0.5)
plt.axhline(y=0, color='r', linestyle='--')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title(f'Residuals Plot for {best_model_name}')
plt.show()
```

- b) **Predicted vs Actual plot:** This plot compares predicted values to actual values, giving a visual representation of how well the model performs.

```
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred, alpha=0.5)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], 'r--')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title(f'Predicted vs Actual Values for {best_model_name}')
plt.show()
```

- c) **Learning curves:** Learning curves show how the model's performance changes with varying amounts of training data.

```
from sklearn.model_selection import learning_curve

train_sizes, train_scores, test_scores = learning_curve(best_model, X, y,
cv=5, scoring='neg_mean_squared_error')
```

```
plt.figure(figsize=(10, 6))
plt.plot(train_sizes,      -train_scores.mean(axis=1),      label='Training
error')
plt.plot(train_sizes,      -test_scores.mean(axis=1),      label='Validation
error')
plt.xlabel('Training Set Size')
plt.ylabel('Mean Squared Error')
plt.title(f'Learning Curves for {best_model_name}')
plt.legend()
plt.show()
```

d) Feature Distribution Plot:

```
df[['comment_length']].plot(kind='hist', bins=30, alpha=0.7)
plt.xlabel('Comment Length')
plt.title('Feature Distribution')
plt.show()
```

e) Model Performance Comparison Box Plot: A box plot of cross-validation scores can help visualize the variability in model performance

```
from sklearn.model_selection import cross_val_score

cv_results = {}
for name, model in models.items():
    scores = cross_val_score(model, X, y, cv=5, scoring='neg_mean_squared_error')
    cv_results[name] = -scores

cv_results_df = pd.DataFrame(cv_results)
cv_results_df.plot.box(figsize=(10, 6))
plt.title('Model Performance Comparison (Cross-Validation)')
plt.ylabel('Mean Squared Error')
plt.show()
```

f) Error Distribution Plot: Visualize the distribution of errors (residuals) to understand the error characteristics

```
errors = y_test - y_pred

plt.figure(figsize=(10, 6))
sns.histplot(errors, kde=True)
plt.xlabel('Error')
plt.ylabel('Frequency')
plt.title(f'Error Distribution for {best_model_name}')
plt.show()
```

Step 5: Co-occurrence network: Identify clusters of related topics or terms.

```
from sklearn.feature_extraction.text import CountVectorizer
from networkx import Graph
```

```

# Example comments
comments = ['Love the new product', 'Great service and product', 'Amazing
product quality']

# Create a count matrix
vectorizer = CountVectorizer()
X = vectorizer.fit_transform(comments)
terms = vectorizer.get_feature_names_out()

# Create a co-occurrence matrix
co_occurrence = (X.T @ X).toarray()
np.fill_diagonal(co_occurrence, 0)

# Build the network
G = Graph()
for i in range(len(terms)):
    G.add_node(terms[i])

for i in range(len(terms)):
    for j in range(i+1, len(terms)):
        if co_occurrence[i, j] > 0:
            G.add_edge(terms[i], terms[j], weight=co_occurrence[i, j])

# Draw the network
pos = nx.spring_layout(G)
nx.draw(G, pos, with_labels=True, node_size=3000, node_color='lightgreen',
font_size=10, font_weight='bold')
edge_labels = nx.get_edge_attributes(G, 'weight')
nx.draw_networkx_edge_labels(G, pos, edge_labels=edge_labels)
plt.title('Co-occurrence Network')
plt.show()

```